




Do this week's survey!!
yellkey.com/set

Lecture 27

Hash Tables II

CS61B, Spring 2026 @ UC Berkeley

Josh Hug and Kay Ousterhout



Hash Table Recap

Lecture 27, CS61B, Spring 2026

Hash Table Recap

Visualization for Some Basic Cases

hashCode and Equals

- Why Custom Hash Functions?
- contains
- Duplicate Values

The Danger of Mutable Keys



- Mutable vs. Immutable Types
- Mutable Hash Table Keys

A Peek into Java HashSets

Goal: Fast Set and Map Implementation

Problem 1:


How do we translate data to an array index?



Goal: Fast Set and Map Implementation


Problem 1:

How do we translate data to an array index?



Problem 2:

What if items have the same index?




Goal: Fast Set and Map Implementation

Problem 1:


How do we translate data to an array index?

Data is converted to a **hash code**;
mod to get a valid index



Problem 2:

What if items have the same index?




Goal: Fast Set and Map Implementation

Problem 1:

How do we translate data to an array index?


Data is converted to a **hash code**;
mod to get a valid index




Problem 2:

What if items have the same index?

Each bucket is a **separate chain** of items
Resize as total items grows



	contains(x)	add(x)
Bushy Binary Search Trees	?	?
Separate Chaining Hash Table With Resizing	?	?





Goal: Fast Set and Map Implementation

	contains(x)	add(x)
Bushy Binary Search Trees	$\Theta(\log N)$	$\Theta(\log N)$
Separate Chaining Hash Table With Resizing	$\Theta(1)^\dagger$	$\Theta(1)^*\dagger$

*: Indicates “on average”.

†: Assuming items are evenly spread.




Goal: Fast Set and Map Implementation

Rest of today

Problem 1:

How do we translate data to an array index?


Data is converted to a **hash code**;
mod to get a valid index



Problem 2:


What if items have the same index?

Each bucket is a **separate chain** of items
Resize as total items grows



Hash tables:

- *Data* is converted into a hash code, the **hash code** is then **reduced** to a valid *index*.
- *Data* is then stored in a bucket corresponding to that *index*.
 - Each bucket is a “separate chain” of items.
- Resize when load factor N/M exceeds some constant.
- If items are spread out nicely, you get $\Theta(1)$ average runtime.



Hash based set ops

	contains(x)	add(x)
Bushy BSTs	$\Theta(\log N)$	$\Theta(\log N)$
Separate Chaining Hash Table With No Resizing	$\Theta(N)$	$\Theta(N)$
... With Reasizing	$\Theta(1)^\dagger$	$\Theta(1)^{* \dagger}$

*: Indicates “on average”.

†: Assuming items are evenly spread.

Visualization for Some Basic Cases

Lecture 27, CS61B, Spring 2026

Hash Table Recap

Visualization for Some Basic Cases

hashCode and Equals

- Why Custom Hash Functions?
- contains
- Duplicate Values

The Danger of Mutable Keys

- Mutable vs. Immutable Types
- Mutable Hash Table Keys

A Peek into Java HashSets

Let's play around with a hash table visualizer.

Goal, get a deeper understanding of:

- How hash codes affect the distribution of items.
- The interaction between equals and hashCode.
- Why hash tables are fast even though they use linked lists.

The objects we're inserting into the HashTable are of type ColoredNumber. Each has 2 attributes:

```
private int num;  
private Color color;
```

Let's see what happens when we insert ColoredNumbers 0 through 19 into a hash table with 6 buckets.

What do you notice about the distribution of items?

Why do you we get this distribution?

0:	2	19			
1:	0	12	13	14	15
2:	3	8	9		
3:	1	7	11	18	
4:	10	16			
5:	4	5	6	17	


Distribution of Items (Your Answer)

What do you notice about the distribution of items?

- Different colored numbers end up in the same bin.
- Distribution is roughly even across bins.
- Looks kinda like a bell curve...sometimes.
- Sequential items in the same bin...sometimes.

Why do you think we get this distribution?

- The default hashCode in Java uses the memory address of the object.



Creating a Custom hashCode Function


Suppose we create a hashCode function that returns 0.

```
@Override  
public int hashCode() {  
    return 0;  
}
```

What distribution do we expect?

- In other words, how will the figure to the right change?

Let's try it out.



The Zero HashCode Function

The resulting distribution is to the right.

0: 0 1 2 3 4 5 6 7 8

1:

2:

3:

4:

5:

Designing a Hash Function

What hash function will result in the distribution to the right?

```
private int num;  
private Color color;
```

0:	0	6	12	18
1:	1	7	13	19
2:	2	8	14	
3:	3	9	15	
4:	4	10	16	
5:	5	11	17	

Designing a Hash Function

What hash function will result in the distribution to the right?

```
private int num;  
private Color color;
```

Your answer?

•

Let's check it.

0:	0	6	12	18
1:	1	7	13	19
2:	2	8	14	
3:	3	9	15	
4:	4	10	16	
5:	5	11	17	

We can define whatever hash function we want:

- Leading digit.
- Sum of the digits.
- Length of the number.
- Something else?

Why Custom Hash Functions?

Lecture 27, CS61B, Spring 2026

Hash Table Recap

Visualization for Some Basic Cases

hashCode and Equals

- **Why Custom Hash Functions?**
- contains
- Duplicate Values

The Danger of Mutable Keys

- Mutable vs. Immutable Types
- Mutable Hash Table Keys

A Peek into Java HashSets

We mentioned that the goal of a hash function is to try to spread items out evenly.

- No spread: Returning 0.
- Bad spread: Returning sum of the digits.
- Good spread: Returning num.

We mentioned that the goal of a hash function is to try to spread items out evenly.

- No spread: Returning 0.
- Bad spread: Returning sum of the digits.
- Good spread: Returning num.

What do you think about the spread of the default hashCode, which returns the memory address?

- A. No spread.
- B. Bad spread.
- C. Good spread.

We mentioned that the goal of a hash function is to try to spread items out evenly.

- No spread: Returning 0.
- Bad spread: Returning sum of the digits.
- Good spread: Returning num.

What do you think about the spread of the default hashCode, which returns the memory address?

- A. No spread.
- B. Bad spread.
- C. **Good spread: The memory address is effectively random, so items should be evenly distributed.**

Hard Question

If the default hashCode achieves good spread, why do we even bother to create custom hash functions?

0:	1	4			
1:	2	14	15	16	17
2:	5	10	11		
3:	3	9	13		
4:	0	12	18		
5:	6	7	8	19	

contains

Lecture 27, CS61B, Spring 2026

Visualization for Some Basic Cases

hashCode and Equals

- Why Custom Hash Functions?
- **contains**
- Duplicate Values

The Danger of Mutable Keys

- Mutable vs. Immutable Types
- Mutable Hash Table Keys

A Peek into Java HashSets

Assume we use the default hash function (the object's memory address).

```
HashSet<ColoredNumber> hs = new HashSet<>();  
ColoredNumber x = new ColoredNumber(4);  
hs.add(x);  
IO.println(hs.contains(x));  
  
ColoredNumber y = new ColoredNumber(4);  
IO.println(hs.contains(y));
```



What's the return value of each call to contains()?

Assume we use the default hash function (the object's memory address).

```
HashSet<ColoredNumber> hs = new HashSet<>();  
ColoredNumber x = new ColoredNumber(4);  
hs.add(x);  
IO.println(hs.contains(x)); // true  
  
ColoredNumber y = new ColoredNumber(4);  
IO.println(hs.contains(y)); // false
```

What do you think the second contains() call should return? Why?

Are the two ColoredNumbers equal?

```
IO.println(x.equals(y)); // false
```

The equals Method for a ColoredNumber

Suppose the equals method for ColoredNumber is as below, i.e. two ColoredNumbers are equal if they have the same num.

- General principle: equal objects should behave identically

```
@Override
public boolean equals(Object o) {
    if (o instanceof ColoredNumber otherCn) {
        return this.num == otherCn.num;
    }
    return false;
}
```

Now that we've written a new equals() method, what's returned by each contains call?

```
@Override
public boolean equals(Object o) {
    if (o instanceof ColoredNumber otherCn) {
        return this.num == otherCn.num;
    }
    return false;
}
```



```
HashSet<ColoredNumber> hs = new HashSet<>();
ColoredNumber x = new ColoredNumber(4);
hs.add(x);
IO.println(hs.contains(x));

ColoredNumber y = new ColoredNumber(4);
IO.println(hs.contains(y));
```




Contains, take 2?

```
HashSet<ColoredNumber> hs = new HashSet<>();  
ColoredNumber x = new ColoredNumber(4);  
hs.add(x);  
IO.println(hs.contains(x));  
  
ColoredNumber y = new ColoredNumber(4);  
IO.println(hs.contains(y));
```




Contains, take 2?

```
HashSet<ColoredNumber> hs = new HashSet<>();  
ColoredNumber x = new ColoredNumber(4);  
hs.add(x);  
IO.println(hs.contains(x));  
  
ColoredNumber y = new ColoredNumber(4);  
IO.println(hs.contains(y)); // false
```



Could contains() have returned true?


```
HashSet<ColoredNumber> hs = new HashSet<>();  
ColoredNumber x = new ColoredNumber(4);  
hs.add(x);  
IO.println(hs.contains(x));  
  
ColoredNumber z = new ColoredNumber(4);  
IO.println(hs.contains(z)); // true
```



contains() with the default Hash Code

```
HashSet<ColoredNumber> hs = new HashSet<>();  
ColoredNumber x = new ColoredNumber(4);  
hs.add(x);  
IO.println(hs.contains(x));  
  
ColoredNumber z = new ColoredNumber(4);  
IO.println(hs.contains(z));
```

contains() returns true
1 / (# buckets)
probability;
otherwise false



Finding an Item Using the Default hashCode

hashCode: Based on memory address.

equals: Based on num.


```
ColoredNumber y = new ColoredNumber(4);  
hs.contains(y); // returns ??
```

There are two ColoredNumber objects with num = 4.

- One of them is in the HashSet.
- One of them was created by the code above.

Each memory address is random.

- Only $1/\#$ buckets chance they % to the same bucket.



Duplicate Values

Lecture 27, CS61B, Spring 2026

Hash Table Recap

Visualization for Some Basic Cases

hashCode and Equals

- Why Custom Hash Functions?
- contains
- **Duplicate Values**

The Danger of Mutable Keys

- Mutable vs. Immutable Types
- Mutable Hash Table Keys

A Peek into Java HashSets

Suppose we have the same equals method (comparing num), but **we do not override hashCode**.

- Result of adding 0 through 19 is shown to the right.

```
public boolean equals(Object o) {  
    ... return this.num == otherCn.num; ...  
}
```

```
ColoredNumber zero = new ColoredNumber(0);  
hs.add(zero); // does another zero appear?
```

Which can happen when we call add(zero)?

- A. We add another 0 to bin zero.
- B. We add another 0 to bin one.
- C. We add another 0 to some other bin.
- D. We do not get a duplicate zero.



0:	2	19			
1:	0	12	13	14	15
2:	3	8	9		
3:	1	7	11	18	
4:	10	16			
5:	4	5	6	17	

Overriding Equals but Not hashCode


Suppose we have the following equals method, but **we do not override hashCode**.

- Result of adding 0 through 19 is shown to the right.

```
@Override
public boolean equals(Object o) {
    if (o instanceof ColoredNumber otherCn) {
        return this.num == otherCn.num;
    }
    return false;
}
```

Which of the following can happen if we add 0 again?

- A. We add another 0 to bin zero.
- B. We add another 0 to bin one.
- C. We add another 0 to some other bin.
- D. We do not get a duplicate zero.



Overriding Equals but Not hashCode

Suppose we have the following equals method, but **we do not override hashCode**.


- Result of adding 0 through 19 is shown to the right.

```
@Override
public boolean equals(Object o) {
    if (o instanceof ColoredNumber otherCn) {
        return this.num == otherCn.num;
    }
    return false;
}
```

Which of the following can happen if we add 0 again?

The new zero ends up in a random bin.

- 5/6ths chance: In bin 0, 2, 3, 4, or 5. Duplicate!
- 1/6 chance: In bin 1, no duplicate! (equals blocks it)



Bottom line: If your class override `equals`, you should also override `hashCode` in a consistent manner.

- If two objects are `equals`, they must always have the same `hashCode`.

If you don't everything breaks:

- Contains can't find objects (unless it gets lucky).
- Add results in duplicates.

Mutable vs. Immutable Types

Lecture 27, CS61B, Spring 2026

Hash Table Recap

Visualization for Some Basic Cases

hashCode and Equals

- Why Custom Hash Functions?
- contains
- Duplicate Values

The Danger of Mutable Keys

- **Mutable vs. Immutable Types**
- Mutable Hash Table Keys

A Peek into Java HashSets

An immutable data type is one for which an instance cannot change in any observable way after instantiation.

Examples:

- Mutable: ArrayDeque, Percolation.
- Immutable: Integer, String, Date.

```
public class Date {  
    public final int month;  
    public final int day;  
    public final int year;  
    private boolean contrived = true;  
    public Date(int m, int d, int y) {  
        month = m; day = d; year = y;  
    }  
}
```

The ***final*** keyword will help the compiler ensure immutability.

- final variable means you may assign a value once (either in constructor of class or in initializer), but after it can never change.
- Final is neither sufficient nor necessary for a class to be immutable.

Which of the Classes Below Are Immutable?

Immutable: an instance cannot change in any observable way after instantiation.

```
public class Pebble {  
    public int weight;  
    public Pebble() {  
        weight = 1;  
    }  
}
```

```
public class Rock {  
    public final int weight;  
    public Rock (int w) {  
        weight = w;  
    }  
}
```

```
public class RocksBox {  
    public final Rock[] rocks;  
    public RocksBox (Rock[] rox) {  
        rocks = rox;  
    }  
}
```

```
public class SecretRocksBox {  
    private Rock[] rocks;  
    public SecretRocksBox(Rock[] rox) {  
        rocks = rox;  
    }  
}
```

Which of the Classes Below Are Immutable?

Immutable: an instance cannot change in any observable way after instantiation.

```
public class Pebble {  
    public int weight;  
    public Pebble() {  
        weight = 1;  
    }  
}
```

Example mutation:

```
Pebble p = new Pebble();  
p.weight = 2;
```

Which of the Classes Below Are Immutable?

Immutable: an instance cannot change in any observable way after instantiation.

No mutation possible.

```
public class Rock {  
    public final int weight;  
    public Rock (int w) {  
        weight = w;  
    }  
}
```

Unless you use the special “Reflections” library which lets you disobey access modifiers.

Which of the Classes Below Are Immutable?

Immutable: an instance cannot change in any observable way after instantiation.

Example mutation:

```
Rock r1 = new Rock(10);  
Rock r2 = new Rock(20);  
Rock[] rox = {r1, r2};  
RocksBox rb = new RocksBox(rox);  
rb.rocks[1] = null;
```

```
public class RocksBox {  
    public final Rock[] rocks;  
    public RocksBox (Rock[] rox) {  
        rocks = rox;  
    }  
}
```

Which of the Classes Below Are Immutable?

Immutable: an instance cannot change in any observable way after instantiation.

Example mutation:

```
Rock r1 = new Rock(10);  
Rock r2 = new Rock(20);  
Rock[] rox = {r1, r2};  
SecretRocksBox rb = new SecretRocksBox(rox);  
rox[0] = new Rock(-999);
```

```
public class SecretRocksBox {  
    private Rock[] rocks;  
    public SecretRocksBox(Rock[] rox) {  
        rocks = rox;  
    }  
}
```

Which of the Classes Below Are Immutable?

Immutable: an instance cannot change in any observable way after instantiation.

```
public class Pebble {  
    public int weight;  
    public Pebble() {  
        weight = 1;  
    }  
}
```

```
public class Rock {  
    public final int weight;  
    public Rock (int w) {  
        weight = w;  
    }  
}
```

```
public class RocksBox {  
    public final Rock[] rocks;  
    public RocksBox (Rock[] rox) {  
        rocks = rox;  
    }  
}
```

```
public class SecretRocksBox {  
    private Rock[] rocks;  
    public SecretRocksBox(Rock[] rox) {  
        rocks = rox;  
    }  
}
```


How Would We Make SecretRocksBox Immutable?

Immutable: an instance cannot change in any observable way after instantiation.

Example mutation:

```
Rock r1 = new Rock(10);  
Rock r2 = new Rock(20);  
Rock[] rox = {r1, r2};  
SecretRocksBox rb = new SecretRocksBox(rox);  
rox[0] = new Rock(-999);
```

```
public class SecretRocksBox {  
    private Rock[] rocks;  
    public SecretRocksBox(Rock[] rox) {  
        rocks = rox;  
    }  
}
```

How Would We Make SecretRocksBox Immutable?

To make SecretRocksBox immutable, we can make our own copy of the array.

- Example mutation fails!

```
Rock r1 = new Rock(10);
Rock r2 = new Rock(20);
Rock[] rox = {r1, r2};
SecretRocksBox rb = new SecretRocksBox(rox);
rox[0] = new Rock(-999);
```


```
public class SecretRocks {
    private Rock[] rocks;
    public SecretRocks(Rock[] rox) {
        rocks = new Rock[rox.length];
        System.arraycopy(rox, 0,
                           rocks, 0,
                           rox.length);
    }
}
```

Advantage: Less to think about: Avoids bugs and makes debugging easier.

- Analogy: Immutable classes have some buttons you can press / windows you can look inside. Results are ALWAYS the same, no matter what.

Disadvantage: Must create a new object anytime anything changes.

- Example: String concatenation is slow!



Mutable Hash Table Keys

Lecture 27, CS61B, Spring 2026

Hash Table Recap

Visualization for Some Basic Cases

hashCode and Equals

- Why Custom Hash Functions?
- contains
- Duplicate Values

The Danger of Mutable Keys

- Mutable vs. Immutable Types
- **Mutable Hash Table Keys**

A Peek into Java HashSets

In principle, we can create a `HashSet<List>`.

Weird stuff happens if:

- We insert a `List` into a `HashSet`.
- Later mutate that `List`.

Example: hashCode

Consider an ArrayList equal to [0, 1].

- Such a list has hashCode 962 (can compute using code shown).


```
List<Integer> items = new ArrayList<>();  
items.add(0);  
items.add(1);  
System.out.println(items.hashCode());
```

Example

Consider an ArrayList equal to [0, 1].

- Such a list has hashCode 962 (can compute using code shown).

If we add this list to a HashSet with 4 buckets, it lands in bucket 2 ($962 \% 4 = 2$).




```
List<Integer> items = new ArrayList<>();
items.add(0);
items.add(1);
System.out.println(items.hashCode());

HashSet<List<Integer>> hs = new HashSet<>();
hs.add(items);
```

Example

First, we added `[0, 1]`, which had hashCode 962, and landed in bucket 2.

- Suppose we now add the list `[2, 3]`. This list has hashCode 1026, which also lands in bucket 2.




```
List<Integer> items = new ArrayList<>();
items.add(0);
items.add(1);
HashSet<List<Integer>> hs = new HashSet<>();
hs.add(items);
hs.add(List.of(2, 3));
```


Example

Suppose we add [0, 1], then [2, 3].

Now suppose we add the number 7 to items.



```
List<Integer> items = new ArrayList<>();
items.add(0);
items.add(1);
HashSet<List<Integer>> hs = new HashSet<>();
hs.add(items);
hs.add(List.of(2, 3));
items.add(7);
```


Example

Suppose we add [0, 1], then [2, 3].

- Now suppose we add the number 7 to items.

The hashCode of a list with [0, 1, 7] is 29829.

- What could go wrong?
- What method call will fail?



```
List<Integer> items = new ArrayList<>();
items.add(0);
items.add(1);
HashSet<List<Integer>> hs = new HashSet<>();
hs.add(items);
hs.add(List.of(2, 3));
items.add(7);
```


Example

Suppose we add [0, 1], then [2, 3].

- Now suppose we add the number 7 to items.

The hashCode of a list with [0, 1, 7] is 29829.

- What could go wrong?
- What method call will fail?
 - contains(items)



```
List<Integer> items = new ArrayList<>();
items.add(0);
items.add(1);
HashSet<List<Integer>> hs = new HashSet<>();
hs.add(items);
hs.add(List.of(2, 3));
items.add(7);
System.out.println(hs.contains(items));
```


Example

Suppose we add [0, 1], then [2, 3].

- Now suppose we add the number 7 to items.
- The hashCode of a list with [0, 1, 7] is 29829.

If we call contains(items), we have a problem.

- hashCode of items is $29829 \% 4 = 1$.
- Hash table looks in bucket 1, empty!




```
List<Integer> items = new ArrayList<>();
items.add(0);
items.add(1);
HashSet<List<Integer>> hs = new HashSet<>();
hs.add(items);
hs.add(List.of(2, 3));
items.add(7);
System.out.println(hs.contains(items));
```

Don't Mutate Keys

Bottom line: Never mutate an Object being used as a key.

- Incorrect results, item gets lost.



```
List<Integer> items = new ArrayList<>();
items.add(0);
items.add(1);
HashSet<List<Integer>> hs = new HashSet<>();
hs.add(items);
hs.add(List.of(2, 3));
items.add(7);
System.out.println(hs.contains(items));
```

A Peek into Java HashSets

Lecture 27, CS61B, Spring 2026

Hash Table Recap

Visualization for Some Basic Cases

hashCode and Equals

- Why Custom Hash Functions?
- contains
- Duplicate Values

The Danger of Mutable Keys

- Mutable vs. Immutable Types
- Mutable Hash Table Keys

A Peek into Java HashSets

We can look at the code that implements the HashSet in Java:

- <https://github.com/openjdk/jdk/blob/master/src/java.base/share/classes/java/util/HashSet.java>


It simply delegates all of its work to a `HashMap<K, Object>` and ignores the value.

We can then look at the code that implements the HashMap in Java:

- <https://github.com/openjdk/jdk/blob/master/src/java.base/share/classes/java/util/HashMap.java>

Reading the code, we can see that:

- Hash table starts at size 16, then doubles every time N exceeds load factor which defaults to 0.75.
- The reduce function is a bit complicated. Come ask me at OH if you're curious.



We can print out the size of the array holding the buckets (requires using the reflections library):

- N/M is never more than 0.75.

Resize occurred, $N = 13$, hash table array size = 32

Resize occurred, $N = 25$, hash table array size = 64

Resize occurred, $N = 49$, hash table array size = 128

Resize occurred, $N = 97$, hash table array size = 256

Resize occurred, $N = 193$, hash table array size = 512

Resize occurred, $N = 385$, hash table array size = 1024

Resize occurred, $N = 769$, hash table array size = 2048

Resize occurred, $N = 1537$, hash table array size = 4096

Visualizing a Hash Table With Load Factor 0.75.

Let's run a simulation to see what happens if the load factor is kept to 0.75 or less.

We can simulate what happens if the load factor is kept to 0.75 or less.

Longest bucket in the simulation ends up around ~ 5 .

- In CS70, you'll derive a precise mathematical characterization of the length of the longest bucket.
 - This is the so called “balls into bins” problem.
 - https://en.wikipedia.org/wiki/Balls_into_bins_problem
 - Can use this math to show the worst case runtime is a bit worse than constant.

If we ctrl-F for “red-black” we find that that if a bin gets too full, it is converted into a red-black tree!

- “This map usually acts as a binned (bucketed) hash table, but when bins get too large, they are transformed into bins of `TreeNode`s, each structured similarly to those in `java.util.TreeMap`. Most methods try to use normal bins, but relay to `TreeNode` methods when applicable (simply by checking `instanceof` a node). Bins of `TreeNode`s may be traversed and used like any others, but additionally support faster lookup when overpopulated. However, since the vast majority of bins in normal use are not overpopulated, checking for existence of tree bins may be delayed in the course of table methods.”
- This is well beyond the scope of our course.


“The most useful algorithms are, unfortunately, not always the most beautiful.”

-Josh Hug

Hashbrowns in Cyberspace by Dall-E.

- Josh conjectures that H thing is the hashbrown key.

(wow this is an old slide! Gen ai is a lot better now)



What is the distinction between hash set, hash map, and hash table?

A hash set is an implementation of the Set ADT using the “hash table” as its engine.

A hash map is an implementation of the Map ADT using the “hash table” as its engine.

A “hash table” is a way of storing information, where you have M buckets that store N items. Each item has a “hashCode” that tells you which of M buckets to put that item in.